

## Computer Memory Structure

1. Processor
    - a. CPU
      - i. Super fast
      - ii. Super expensive
      - iii. Tiny capacity
    - b. CPU Cache (L1 - L3 Cache)
      - i. Not as fast as CPU
      - ii. Expensive
      - iii. Small capacity
  2. RAM
    - a. Physical Memory (Random Access Memory)
      - i. Not as fast as CPU cache
      - ii. Priced reasonably
      - iii. Average Capacity
  3. Solid State Drives (Solid State Memory)
    - a. Non-Volatile Flash-Based Memory
      - i. Average Speed
      - ii. Priced reasonable
      - iii. Average capacity
  4. Mechanical Hard Drives (Virtual Memory)
    - a. File-Based Memory
      - i. Slow
      - ii. Cheap
      - iii. Large Capacity
- For fast databases, we want to minimize access to solid state drives and hard drives
  - A 64-bit integer takes up 8 bytes of memory
  - Say you have a 2048 byte block size. To get 1, 64-bit (8 byte) integer out of memory, the DB has to read 2048 bytes.
  - In an AVL node containing a key-value pair of two 64-bit integers, we are using 32 bytes of memory
    - 8 bytes for key, 8 bytes for value, 8 bytes for right child pointer, 8 bytes for left child pointer
    - Therefore, we are not optimizing the amount of memory we have available. We can increase the performance by decreasing the height of the tree
  - Consider the following case: We have a sorted array of 128 integers

- $128 * 8 = 1024$  bytes
- In the worst case, binary search o 128 integers is still much faster than a single additional disk access
- Reading only one block of memory from a hard drive is significantly faster than reading for a second time
- What if we could have two keys per node? Or three or four...?
  - As you add more keys to each node, the tree gets shallower (smaller height) and searching gets faster
  - The main point: in databases, we want to minimize hard drive access
- The ideal structure or databases and indexing are B+ trees
  - Each node can have up to 128, 256, etc values (keys) for indexing.
  - For n-1 keys, we can have n children
  - Searching for a child is still much faster than another disk read
  - This way, we can store more values and keys in much fewer levels
- The goal of indexing structures = minimize the height of the tree