

1. Redis in Machine Learning:

- a. Pipeline: Data Source (AWS) -> Transformations (Apache Spark) -> Inference Store (Stored on Redis) + Training Store (Stored on AWS)
 - i. Question: Why would you add complexity by adding the inference store on Redis?
 - ii. Answer: For latency issues: accessing a key from Redis is much faster than running a SELECT statement. Redis is stored on disk which is faster than accessing from RAM.

2. Port mapping a:b

- a. A = host port on the physical machine
- b. B = port inside the container
- c. Maximum port number = 65535
- d. At the operating system level, http maps to port 80, https maps to port 443, ssh maps to port 22, ftp maps to 21
- e. Range of ports reserved for system services: 0 - 102